

Chapter 2 (DRAFT - 3/1/16)

Hugh Zabriskie

7 March 2016

1 Motivation for the harmonization task

The task explored in this paper is the generation of four-part counterpoint given a chorale melody. J.S. Bach's collection of over 300 chorale harmonizations is used as the dataset because it remains today one of the finest examples of four-part counterpoint. Each harmonization is a complex solution, satisfying a variety of voice leading and cadential constraints, as well as innumerable other conventions of musical style. Chorale harmonization is therefore a fundamentally difficult computational task. The reader might ask whether a rule-based approach is suitable to this task, since after all, isn't harmonization a largely rule-based. Couldn't musical constraints would be computationally encoded and then only harmonizations that satisfy them are generated? However, a purely rule-based approach is rendered impractical by a couple important factors. One is the sheer number of conventions that would need to be encoded, each with a varying level of specificity and precedence with respect to the other encoded conventions. Determining precisely what those conventions would be is an even more difficult problem, since many chorale constraints are vaguely defined or flexible in application. Another reason for avoiding a rule-based approach is that an infinite quantity of harmonizations exist that satisfy any given chorale melody. The goal is not to simply produce satisfying harmonizations, but to *approximate* Bach's harmonization as closely as possible. In

23 contrast, a "learning" model updates its parameters to capture the complex corre-
24 lations and patterns it discovers in the training set of harmonizations. Applying
25 those parameters to a new chorale melody will then yield a predicted harmoniza-
26 tion, and its predictions will directly reflect (and only reflect) the harmonizations
27 the model was trained upon.

28 **2 Literature Review**

29 Hild, Feulner, and Menzel (1992) presented the first effective neural network ap-
30 proach for harmonizing chorales, generation harmonizations on the level of an
31 "improvising organist" (p. 272). The task decomposed into three subtasks, each
32 learned by a neural net. A harmonic skeleton is first created by sweeping through
33 the chorale melody and determining a harmony for each beat, where harmony is
34 represented by a unique figured bass notation. For each time step t , the network
35 takes an input a window of information, including the harmonies chosen in the
36 interval $[t - 3, t - 1]$ and the melody pitches in the interval $[t - 1, t + 1]$. The result-
37 ing harmonic skeleton is passed to another network to generate a chord skeleton,
38 which selects the inner voices based on the input figured bass, and a final network
39 adds ornamental eighth notes to restore passing tones between chords. Although,
40 HARMONET demonstrated strong success, the model used external "chorale con-
41 straints" (ibid., p. 271) in constructing the chord skeleton in order to avoid unfa-
42 vorable chord structures. The models presented in this paper attempt to learn the
43 task of harmonization without any form of manual intervention in the network's
44 learning process.

45
46 Substantial work has been done in the field of music generation using RNNs
47 and LSTMs that demonstrates their ability to learn complicated musical structures
48 and relate temporally distant events, particularly in the realm of melodic gener-

49 ation. Toivainen (1995) developed a neural network that generates a jazz bebop
50 melody over series of chord changes. The network achieves melodic continuity
51 by using a "target-note technique", where the end of a previous melody segment
52 and the present chord are used to predict the next melodic pattern at time $t + 1$,
53 while the following chord at time $t + 2$ is use to optimize of the melodic pattern,
54 thereby smoothing the melodic transitions over chord changes. Eck and Schmid-
55 huber (2002) improved on the results of Mozer (1994), who used RNNs to compose
56 melodies with chordal accompaniment but found the resulting music lacked larger
57 thematic and phrase structure. They attributed Mozer's results to the "vanishing
58 gradients" problem (described briefly in chapter 1) and then trained LSTMs to first
59 generate the chordal structure and use that as a input to the LSTMs that gener-
60 ates a blues-style melody with promising results. Franklin (2006) also used LSTMs
61 (two inter-recurrent networks) to compose jazz melodies over a chord progression,
62 training the networks on a dataset of well-known jazz standard melodies.

63

64 Research specifically on models for the Bach chorales have seen a variety of
65 approaches. Allan and Williams (2005) generated a dataset of the chorales by
66 transposing each chorale to C major or C minor and then sampling every quarter
67 note, now available at [http://www-etud.iro.umontreal.ca/~boulanni/](http://www-etud.iro.umontreal.ca/~boulanni/icml2012)
68 [icml2012](http://www-etud.iro.umontreal.ca/~boulanni/icml2012). They trained Hidden Markov Models (HMMs) on the data to model
69 chorale harmonization, creating a probabilistic framework for deciding the most
70 likely choice of alto, tenor, and bass notes to complement the melody in each time
71 frame. More recent work has focused on music *generation* models rather than com-
72 pletive models. Boulanger-Lewandowski et. al. (2012) used this version of the
73 dataset, along with multiple other corpuses, to make a comprehensive survey of
74 music generation models that performed next-step prediction for all voices. Based
75 on log-likelihood and overall accuracy on the test data, a specific flavor of RNNs
76 was found to be most effective on all corpuses used in the study. Other studies

77 have sought to “re-construct” chorales using more sophisticated neural models
78 (Liu, 2014), including a large-scale survey of polyphonic music modeling that eval-
79 uated the performance of eight LSTM variants Greff et al. (2015). Both previous
80 papers mentioned highlight the use of RNNs and LSTMs as effective models be-
81 cause of their ability to learn temporal dependencies within polyphonic music (i.e.
82 how distant musical events can be related). Notably, the latter found the “vanilla”,
83 unmodified LSTM equally effective as the other variants (ibid., p. 7), so this archi-
84 tecture was chosen for this study.

85 My approach consists of applying a variety of non-neural and neural models
86 to the task of chorale harmonization. Recent harmonization and generative mod-
87 els (Allan and Williams, 2005; Kaliakatsos-Papakostas and Cambouropoulos, 2014;
88 Greff et al., 2015) relied on a dataset that includes only pitch information about
89 each time step, and the objective was in all cases next step pitch prediction. My
90 goal was to extract a new set of features from Bach’s chorales and examine how
91 those features improved or worsened model performance. Moreover, the objective
92 was to learn a series of musical processes that decide the harmony at each time
93 step that include both harmonic analysis and note selection. I chose this approach
94 to mimic the decisions a musician would make today in harmonizing a chorale.

95 **2.1 Harmonization tasks**

96 In order to evaluate a model’s ability to learn different aspects of the harmoniza-
97 tion process the decision process was broken into 4 sequential subtasks. For each
98 melody note in each chorale, the following target outputs were extracted.

- 99 1. *Roman numeral*. The model first decided the general harmonic function by
100 selecting a Roman numeral per Roman numeral analysis. This decision holds
101 the most importance since it decides what notes are allowed for the alto and
102 tenor voices.

103 2. *Inversion*. The chord inversion provides additional harmonic information
104 since, for example, a I^6 harmony has different implications about future har-
105 monies than a I chord in root position. The inversion also implies the pitch
106 class of the bass voice.

107 3. *Alto*. The alto voice is then selected as the first inner voice. There is no reason
108 for selecting the alto voice before the tenor voice since both function similarly
109 as inner voices that support the harmony decided in the previous subtasks.

110 4. *Tenor*. Selection of the tenor voice completes the harmonization.

111 Mathematically, the algorithm should take as input the chorale, represented a
112 sequence of notes, and it should output a corresponding sequence of 3-voice chords
113 that represent the alto, tenor, and bass voices. We will say we have m data points
114 with n features, and Y output classes.

- 115 • $\mathbf{X} \in V^{m \times n}$ is our input data
- 116 • $\mathbf{Y} \in Y^m$ is our output data for a given subtask
- 117 • The objective is to approximate the complex function $f : \mathbf{X} \rightarrow \mathbf{Y}$

118 In order to generate the correct Roman numeral analysis for the first two sub-
119 tasks, I relied on a combination of MUSIC21's `roman` module for initial analysis
120 followed by substantial manual correction due some incomplete functionality in
121 the analysis module. The alto and tenor voices were extracted as MIDI note values.


122

123 3 Methods

124 3.1 Important resources

125 We gathered chorale in MusicXML format through the MUSIC21 `corpus` module,
126 which originally obtained the chorales from Greentree (2005). MUSIC21 is a toolkit

Figure 1: **Sequence of harmonization subtasks.**



for computational musicology maintained by Professor Michael Scott Cuthbert at MIT (Cuthbert and Ariza, 2010), and its library is built upon a comprehensive array of objects and methods that represent the fundamental musical components introduced in Chapter 1. For our purposes, MUSIC21 was used extensively to transform musical scores into symbolic representations and extracting musical features from those scores for machine learning. It was also used to generate realizations of predicted chorale harmonizations as actual scores, which are provided in Chapter 3.

In order to construct and train neural networks, the scientific computing framework Torch was used, which is written in the Lua programming language. The `rnn` library, which is designed to extend the capabilities of the `nn` module in Torch, was used to implement recurrent neural networks like RNNs and LSTM networks (Léonard, Waghmare, and Wang, 2015). Non-neural models were selected from the fully implemented classification algorithms in the machine learning library `scikit-learn` (Pedregosa et al., 2011).

3.2 Data

The dataset consists of 326 4-voice chorales gathered from the MUSIC21 library. Once collected, some manual cleaning was then performed to correct mistakes in

145 the musicXML format related to key signatures and significant mistakes in notation
146 (based on a visual comparison with the Riemenschneider edition of the Chorales).
147 Chorales with significant periods of rest were removed from the dataset because 3-
148 voice harmonies have ambiguous implications in the 4-voice model that the chorales
149 conventionally observe. Next, the chorales were “quantized” to create strictly
150 chordal progressions of 4-voice harmony. Like modern church hymns, Bach’s chorales
151 are uniformly structured as chordal progressions, with a consistent beat-long rate
152 of harmonic transition. Therefore, the process of quantizing each chorale into a
153 series of discrete and uniform time steps, each the length of a beat, could be ac-
154 complished without damaging the underlying harmonic progression. Eighth notes
155 were removed from the voices to eliminate additional rhythmic complexity. Eighth
156 notes facilitate voice leading in Bach’s harmonizations, but rarely function as an
157 essential part of the harmony, making their removal a reasonable design decision.
158 This process of quantizing the chorale into quarter-note samples has been found to
159 be effective in several other studies (Hild, Feulner, and Menzel, 1992; Madsen and
160 Jorgensen, 2002; Kaliakatsos-Papakostas and Cambouropoulos, 2014).

161 **3.3 Feature Extraction**

162 In order to learn each of these subtasks, the following information was extracted
163 for each melody note in each chorale.

- 164 1. The number of sharps in the key signature. Flats were given negative values.
- 165 2. The mode (i.e. major or minor) of the chorale.
- 166 3. The time signature of the chorale.
- 167 4. Beat strength, or metrical accent. A 4/4 measure would be assigned the fol-
168 lowing pattern: [1.0, 0.25, 0.5, 0.25]
- 169 5. The presence of a fermata - a binary feature indicating a cadence.

- 170 6. Number of beats until the next fermata.
- 171 7. Number of measures until the end of the chorale.
- 172 8. The melody pitch, encoded as a MIDI value. A search across all chorales
173 indicated that each voice had a well-defined pitch range, verified by Madsen
174 and Jorgensen (2002):
- 175 • *soprano*: [60, 81]
 - 176 • *alto*: [53, 74]
 - 177 • *tenor*: [48, 69]
 - 178 • *bass*: [36, 64]
- 179 9. The interval to the previous melody note.
- 180 10. The interval to the next melody note.
- 181 11. The Roman numeral for the previous time step.
- 182 12. The inversion for the previous time step.

183 Features 11 and 12 were used only for Oracle experiments and demonstrating
184 the potential of RNNs, discussed later in Chapter 3.

185

186 The Python script used to preprocess the chorales (see `wrangle.py` in Appendix
187 B)¹ extracts the above features from each time step of each chorale, and stores the
188 generated training and tests in an HDF5 file. For baseline models, the extracted
189 data was fed as input to another Python script that performed classified learning
190 using models from the `SCIKIT-LEARN` module. For neural models, the data was fed
191 into Lua scripts that use the scientific computing framework `Torch` to construct
192 models and training on supervised classification tasks.

¹There will be an Appendix B that contains the most important code for the project.

4 Baseline models

Due to unique classification problems chosen for this paper, previous computational research on the chorales does not provide any adequate baseline models. Baseline models are important because they provide a basis for comparison when evaluating other models. As the complexity of the model changes and the features are added or removed, it is important to have a metric to compare against to see how those changes improved or worsened the results. In classification problems, a crude baseline model can be achieved by choosing the class with the most observations and using that class as the result for all predictions. TABLE 2 lists the baseline frequencies for the most common classes for each subtask. I also trained three other classifiers to get a sense of the complexity of each subtask. These classifiers are described below:

Multiclass logistic regression was introduced in Chapter 1 as a generalization of the binary classification system of logistic regression. Despite its name, this regression is a linear model. The objective is to minimize the following cost function, given the learned parameters θ .

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K \{y^{(i)} = k\} \log \left(\frac{\exp(a_{ik})}{\sum_j \exp(a_{ij})} \right)$$

Where m is the number of examples, k is the number of classes, and a_{ik} is the "activation" function $\theta^{(k)T} x^{(i)}$, denoted for the i th example and the k th class.

Multinomial naive Bayes generalizes the naive Bayes algorithm for multi-class data, and it makes the "naive" assumption of independence between every pair of input features. The assumption therefore states that, given the input vector x and the class $c \in [1, K]$

$$P(x|c) = P(x_1, x_2, \dots, x_n|c) = \prod_{i=1}^n P(x_i|c)$$

And based on that assumption, this baseline classifier the predicted output class is decided by

$$P(c|\mathbf{x}) = \frac{P(\mathbf{x}|c)P(c)}{\sum_j P(\mathbf{x}|C_j)P(C_j)} = \frac{\prod_i P(x_i|c)P(c)}{\sum_j \prod_i P(x_i|C_j)P(C_j)}$$

196

197

198 *Random forests* are a powerful supervised learning technique that involves clas-
 199 sification based on a the majority vote of a series of decision trees. Each tree is
 200 initialized with data from a random subset of features and then is trained on the
 201 data by sampling with replacement. This randomness is known to be highly effec-
 202 tive in preventing overfitting on training data, and random forests generalize well
 203 on weaker datasets where one or more training examples do not strongly suggest
 204 differences between classes (Breiman, 2001, p. 18).

205 The results for each baseline model are described in Table 1.

206

Table 1: **Baseline model test accuracy on harmonization subtasks.**

Classifier	Numeral	Inversion	Alto	Tenor
Multi-Class Logistic	31.61%	59.76%	37.55%	37.86%
Multinomial Naive Bayes	27.44%	56.66%	35.06%	34.40%
Random Forests	49.29%	61.64%	49.44%	45.43%

Table 2: **Most common class frequency.**

Subtask	Training Set	Test Set
Numeral	19.2%	19.2%
Inversion	55.5%	57.6%
Alto	15.7%	14.6%
Tenor	15.6%	15.5%


4.1 GCT Algorithm

While Roman numeral analysis has been the traditional method for describing harmony in the Chorales, it presents issues for statistical learning. Roman numeral classification mainly depends on the key signature, but also requires the context of the preceding harmonies. For example, a D major chord in a C major chorale might be labelled as a II or V/V depending on whether a modulation to D major had occurred or whether the preceding harmonies indicate that it functions as a secondary dominant. During training, finding two or more inputs that suggest a D major chord but are labeled differently can cause confusion in learning, particularly since in non-recurrent models there is no sense of context about other local harmonies. Roman numeral chord labeling can be further complicated by the presence of non-chord tones, which makes, for example, differentiating IV⁶ and ii⁷ chords computationally difficult.

The general chord type (GCT) representation provides an idiom-independent solution to encoding harmony that assigns a unique encoding to each chord, regardless of context (Cambouropoulos, Kaliakatsos-Papakostas, and Tsougras, 2014). To encode a chordal harmony, the GCT algorithm takes as input the chord to be encoded, a pitch scale that describes the tonality, and a binary "consonance vector" v such that $v[n] = 1$ if an interval of n semitones is considered consonant for $0 \leq n \leq 11$. In this study, I chose $v = [1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0]$. GCT then constructs an ordering of the chord pitches that maximizes the consonant intervals between all pairs of pitches. The remaining notes that create dissonant intervals are labeled as "extensions". The algorithm outputs an encoding of the form [root, [base, extensions]], where root is the pitch class of the chord root relative to the tonic, and the base is the ordering of maximal consonance. I adapted the algorithm to also output the degree of inversion, where 0 represents root position, 1 represents first inversion, and so on. Figure 2 demonstrates an application of the GCT algorithm

235 to a tonal harmonic progression, comparing the Roman numeral analysis with the
 236 GCT encoding. The base $[0, 4, 7]$ encodes a major triad, while $[0, 3, 7, 10]$ encodes a
 237 minor seventh chord.

Figure 2: Example of the GCT and Roman numeral notation.



238 In comparison with a Roman numeral analysis dataset compiled by David Tem-
 239 perley, GCT labeled 92% of the chords accurately, of which about 1/3 of mislabeled
 240 chords were diminished sevenths - excusable because each note in the chord can
 241 function as the root (ibid.). In order to minimize duplicate encodings, I imple-
 242 mented the following policies, partially drawn from suggestions by the original
 243 authors.

- 244 1. For dyads, prefer an interval of a 5th over a 4th, and an interval of a 7th over
 245 a 2nd.
- 246 2. Preference encodings where all intervals are larger than a major 2nd. This
 247 heuristic preferences a minor 7th or a major chord with an added 6th, and
 248 generally more evenly spaced encodings.
- 249 3. If more than one encoding remains, choose randomly.

250 The GCT algorithm was incorporated into a newly generated dataset for the
 251 chorales by replacing the numeral and inversion subtasks in \mathcal{Y} with the new root,
 252 base, and inversion subtasks. The root subtasks establishes the chord root, while
 253 the remaining chord structure is decided by the base. As in Roman numeral anal-
 254 ysis, the inversion implies which chord tone is assigned to the bass. As well, the

new dataset classified the tenor and alto voices by their distance from the tonic pitch, instead of encoding MIDI values directly, in order to make the decision key-independent and reduce the output class space to a maximum of 12 classes (for the 12 chromatic intervals).

The same baseline models were used to evaluate the new subtasks. In both cases, there were 293 chorales in the training set, and 33 chorales in the test set.

Table 3: **Baseline model test accuracy with new GCT subtasks.**


Classifier	Root	Base	Inversion	Alto	Tenor
Multi-Class Logistic	58.29%	56.09%	67.06%	38.22%	36.33%
Multinomial Naive Bayes	48.19%	50.43%	62.66%	34.73%	31.65%
Random Forests	63.71%	59.72%	69.73%	47.22%	46.81%

Table 4: **Majority class frequency (MCCF), with GCT subtasks.**

Subtask	Training Set	Test Set
Root	28.5%	28.0%
Base	49.5%	49.9%
Inversion	63.9%	65.5%
Alto	22.6%	23.7%
Tenor	20.3%	22.0%

Figure 3 provides a visual comparison of baseline model performance across all harmonization subtasks. For all subtasks, a majority of the multinomial models outperformed the MCCF baselines. In particular, Random Forests consistently outperformed the other models, with a 35% increase in accuracy over the MCCF baseline for the root subtask. However, the multinomial models struggled to perform above the MCCF baselines for both the base and inversion subtasks, which appears correlated with a high predominance of a single class in the dataset. This points to an issue with an imbalance class distribution in the data. Of the 133 output classes for the GCT base subtask, the most common class is associated with 50% of all observations in entire dataset (this happens to be the major triad), and

Figure 3: Harmonization subtask accuracy comparison



272 the top 3 most frequent classes account for 77% of all observations. Consequently,
 273 the vast majority of output classes are observed too infrequently in the training
 274 data to be classified accurately in the test data. Imbalanced data is a widely recog-
 275 nized phenomenon in data mining that is comprehensively detailed in Sun, Wong,
 276 and Kamel (2009). The significant advantage achieved using random forests can
 277 potentially be explained by its known effectiveness at classification when one or
 278 more observations is not sufficient to generally distinguish a class (Breiman, 2001,
 279 pg. 18).

References

- Allan, Moray and Christopher KI Williams (2005). "Harmonising chorales by probabilistic inference". In: *Advances in neural information processing systems* 17, pp. 25–32.
- Breiman, Leo (2001). "Random forests". In: *Machine learning* 45.1, pp. 5–32.
- Cambouropoulos, Emilios, Maximos Kaliakatsos-Papakostas, and Costas Tsougras (2014). *An idiom-independent representation of chords for computational music analysis and generation*. Ann Arbor, MI: Michigan Publishing, University of Michigan Library.
- Cuthbert, Michael Scott and Christopher Ariza (2010). "music21: A toolkit for computer-aided musicology and symbolic music data". In: *International Society for Music Information Retrieval*, pp. 637–642.
- Eck, Douglas and Juergen Schmidhuber (2002). "A first look at music composition using lstm recurrent neural networks". In: *Istituto Dalle Molle Di Studi Sull Intelligenza Artificiale*.
- Franklin, Judy A (2006). "Jazz melody generation using recurrent networks and reinforcement learning". In: *International Journal on Artificial Intelligence Tools* 15.04, pp. 623–650.
- Greentree, Margaret, ed. (2005). *Chorales harmonized by J.S. Bach*. Retrieved October 2015 from music21.
- Greff, Klaus et al. (2015). "LSTM: A Search Space Odyssey". In: *arXiv:1503.04069*. URL: <http://adsabs.harvard.edu/abs/2015arXiv150304069G>.
- Hild, Hermann, Johannes Feulner, and Wolfram Menzel (1992). "HARMONET: A neural net for harmonizing chorales in the style of JS Bach". In: *Advances in Neural Information Processing Systems*, pp. 267–274.

305 Kaliakatsos-Papakostas, Maximos and Emilios Cambouropoulos (2014). "Proba-
306 bilistic harmonization with fixed intermediate chord constraints". In: *ICMC—SMC*.
307 Ann Arbor, MI: Michigan Publishing, University of Michigan Library.

308 Léonard, Nicholas, Sagar Waghmare, and Yang Wang (2015). "rnn: Recurrent Li-
309 brary for Torch". In: *arXiv* 1511.07889.

310 Madsen, Soren Tjagvad and Martin Elmer Jorgensen (2002). "Harmonisation of
311 Bach chorales: KBS project report". In: p. 31.

312 Mozer, Michael C (1994). "Neural network music composition by prediction: Ex-
313 ploring the benefits of psychoacoustic constraints and multi-scale processing".
314 In: *Connection Science* 6.2-3, pp. 247–280.

315 Pedregosa, F. et al. (2011). "Scikit-learn: Machine Learning in Python". In: *Journal of*
316 *Machine Learning Research* 12, pp. 2825–2830.

317 Sun, Yanmin, Andrew KC Wong, and Mohamed S Kamel (2009). "Classification of
318 imbalanced data: A review". In: *International Journal of Pattern Recognition and*
319 *Artificial Intelligence* 23.04, pp. 687–719.

320 Toivainen, Petri (1995). "Modeling the Target-Note Technique of Bebop-Style Jazz
321 Improvisation: An Artificial Neural Network Approach". English. In: *Music*
322 *Perception: An Interdisciplinary Journal* 12.4, pp. 399–413. ISSN: 07307829. URL:
323 <http://www.jstor.org/stable/40285674>.